

# Dichotomize and Generalize: PAC-Bayesian Binary Activated Deep Neural Networks

Gaël Letarte<sup>1</sup>, Pascal Germain<sup>2</sup>, Benjamin Guedj<sup>2,3</sup>, François Laviolette<sup>1</sup>

<sup>1</sup> Département d'informatique et de génie logiciel, Université Laval, Québec, Canada

<sup>2</sup> Équipe-projet Modal, Inria Lille - Nord Europe, Villeneuve d'Ascq, France

<sup>3</sup> UCL Centre for Artificial Intelligence, University College London, London, England

The Alan Turing Institute



## Introduction


We present a comprehensive study of multilayer neural networks with binary activation, relying on the PAC-Bayesian theory.

### Contributions

- An end-to-end framework to train a binary activated deep neural network (DNN).
- Nonvacuous PAC-Bayesian generalization bounds for binary activated DNNs.

## Binary Activated Neural Networks

- $L$  fully connected layers
- $d_k$  denotes the number of neurons of the  $k^{\text{th}}$  layer
- $\text{sgn}(a) = 1$  if  $a > 0$  and  $\text{sgn}(a) = -1$  otherwise
- Weights matrices :  $\mathbf{W}_k \in \mathbb{R}^{d_k \times d_{k-1}}$ ,  $\theta = \text{vec}(\{\mathbf{W}_k\}_{k=1}^L) \in \mathbb{R}^D$



### Prediction

$$f_\theta(\mathbf{x}) = \text{sgn}(\mathbf{w}_L \text{sgn}(\mathbf{W}_{L-1} \text{sgn}(\dots \text{sgn}(\mathbf{W}_1 \mathbf{x}))))$$

## PAC-Bayesian Theory

Given a data distribution  $\mathcal{D}$ , a training set  $S = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\} \sim \mathcal{D}^n$ , with  $\mathbf{x}_i \in \mathbb{R}^{d_0}$  and  $y_i \in \{-1, 1\}$ , a loss  $\ell : [-1, 1]^2 \rightarrow [0, 1]$ , a predictor  $f \in \mathcal{F}$  :

$$\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \ell(f(\mathbf{x}), y) \quad \leftarrow \text{generalization loss}$$

$$\hat{\mathcal{L}}_S(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(\mathbf{x}_i), y_i) \quad \leftarrow \text{empirical loss}$$

### PAC-Bayesian Theorem

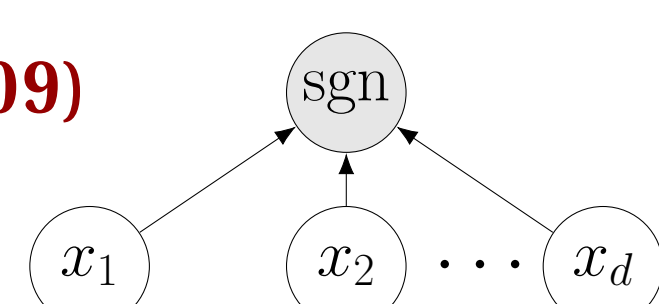
For any prior  $P$  on  $\mathcal{F}$ , with probability  $1 - \delta$  on the choice of  $S \sim \mathcal{D}^n$ , we have for all  $C > 0$ , and any posterior distribution  $Q$  on  $\mathcal{F}$  :

$$\mathbb{E}_{f \sim Q} \mathcal{L}_{\mathcal{D}}(f) \leq \frac{1}{1 - e^{-C}} \left( 1 - \exp \left( -C \mathbb{E}_{f \sim Q} \hat{\mathcal{L}}_S(f) - \frac{1}{n} [\text{KL}(Q \| P) + \ln \frac{2\sqrt{n}}{\delta}] \right) \right)$$

## Linear Classifier


PAC-Bayesian Learning of Linear Classifiers (Germain et al., 2009)

$$f_{\mathbf{w}}(\mathbf{x}) := \text{sgn}(\mathbf{w} \cdot \mathbf{x}), \text{ with } \mathbf{w} \in \mathbb{R}^d$$



### PAC-Bayesian analysis

- Space of all linear classifiers  $\mathcal{F}_d := \{f_v | v \in \mathbb{R}^d\}$
- Gaussian prior  $P_{\mathbf{u}} := \mathcal{N}(\mathbf{u}, I_d)$  over  $\mathcal{F}_d$
- Gaussian posterior  $Q_{\mathbf{w}} := \mathcal{N}(\mathbf{w}, I_d)$  over  $\mathcal{F}_d$
- Predictor  $F_{\mathbf{w}}(\mathbf{x}) := \mathbb{E}_{v \sim Q_{\mathbf{w}}} f_v(\mathbf{x}) = \text{erf}\left(\frac{\mathbf{w} \cdot \mathbf{x}}{\sqrt{2} \|\mathbf{x}\|}\right)$
- Linear loss  $\ell(f_v(\mathbf{x}), y) := \frac{1}{2} - \frac{1}{2} y f_v(\mathbf{x})$



### Bound minimization


$$C n \hat{\mathcal{L}}_S(F_{\mathbf{w}}) + \text{KL}(Q_{\mathbf{w}} \| P_{\mathbf{u}}) = C \frac{1}{2} \sum_{i=1}^n \text{erf}\left(-y_i \frac{\mathbf{w} \cdot \mathbf{x}_i}{\sqrt{2} \|\mathbf{x}_i\|}\right) + \frac{1}{2} \|\mathbf{w} - \mathbf{u}\|^2$$

## Shallow Learning

Posterior  $Q_\theta = \mathcal{N}(\theta, I_D)$ , over the family of all networks  $\mathcal{F}_D = \{f_\theta | \theta \in \mathbb{R}^D\}$ , where

$$f_\theta(\mathbf{x}) = \text{sgn}(\mathbf{w}_2 \cdot \text{sgn}(\mathbf{W}_1 \mathbf{x}))$$

$$\begin{aligned} F_\theta(\mathbf{x}) &= \mathbb{E}_{\theta \sim Q_\theta} f_\theta(\mathbf{x}) \\ &= \int_{\mathbb{R}^{d_1 \times d_0}} Q_1(\mathbf{V}_1) \int_{\mathbb{R}^{d_1}} Q_2(\mathbf{v}_2) \text{sgn}(\mathbf{v}_2 \cdot \text{sgn}(\mathbf{V}_1 \mathbf{x})) d\mathbf{v}_2 d\mathbf{V}_1 \\ &= \sum_{\mathbf{s} \in \{-1, 1\}^{d_1}} \text{erf}\left(\frac{\mathbf{w}_2 \cdot \mathbf{s}}{\sqrt{2} d_1}\right) \int_{\mathbb{R}^{d_1 \times d_0}} \mathbb{1}[\mathbf{s} = \text{sgn}(\mathbf{V}_1 \mathbf{x})] Q_1(\mathbf{V}_1) d\mathbf{V}_1 \\ &= \sum_{\mathbf{s} \in \{-1, 1\}^{d_1}} \underbrace{\text{erf}\left(\frac{\mathbf{w}_2 \cdot \mathbf{s}}{\sqrt{2} d_1}\right)}_{F_{\mathbf{w}_2}(\mathbf{s})} \underbrace{\prod_{i=1}^{d_1} \left[ \frac{1}{2} + \frac{s_i}{2} \text{erf}\left(\frac{\mathbf{w}_1^i \cdot \mathbf{x}}{\sqrt{2} \|\mathbf{x}\|}\right) \right]}_{\Pr(\mathbf{s} | \mathbf{x}, \mathbf{W}_1)} \end{aligned}$$



### PAC-Bayesian bound ingredients

- Empirical loss :  $\hat{\mathcal{L}}_S(F_\theta) = \mathbb{E}_{\theta \sim Q_\theta} \hat{\mathcal{L}}_S(f_\theta) = \frac{1}{n} \sum_{i=1}^n \left[ \frac{1}{2} - \frac{1}{2} y_i F_\theta(\mathbf{x}_i) \right]$
- Complexity term :  $\text{KL}(Q_\theta \| P_\mu) = \frac{1}{2} \|\theta - \mu\|^2$ , with  $\mu \in \mathbb{R}^D$
- Generalization bound :  $\frac{1}{1 - e^{-C}} \left( 1 - \exp \left( -C \hat{\mathcal{L}}_S(F_\theta) - \frac{1}{n} [\text{KL}(Q_\theta \| P_\mu) + \ln \frac{2\sqrt{n}}{\delta}] \right) \right)$

## Visualization

The proposed method can be interpreted as a majority vote of hidden layer representations.

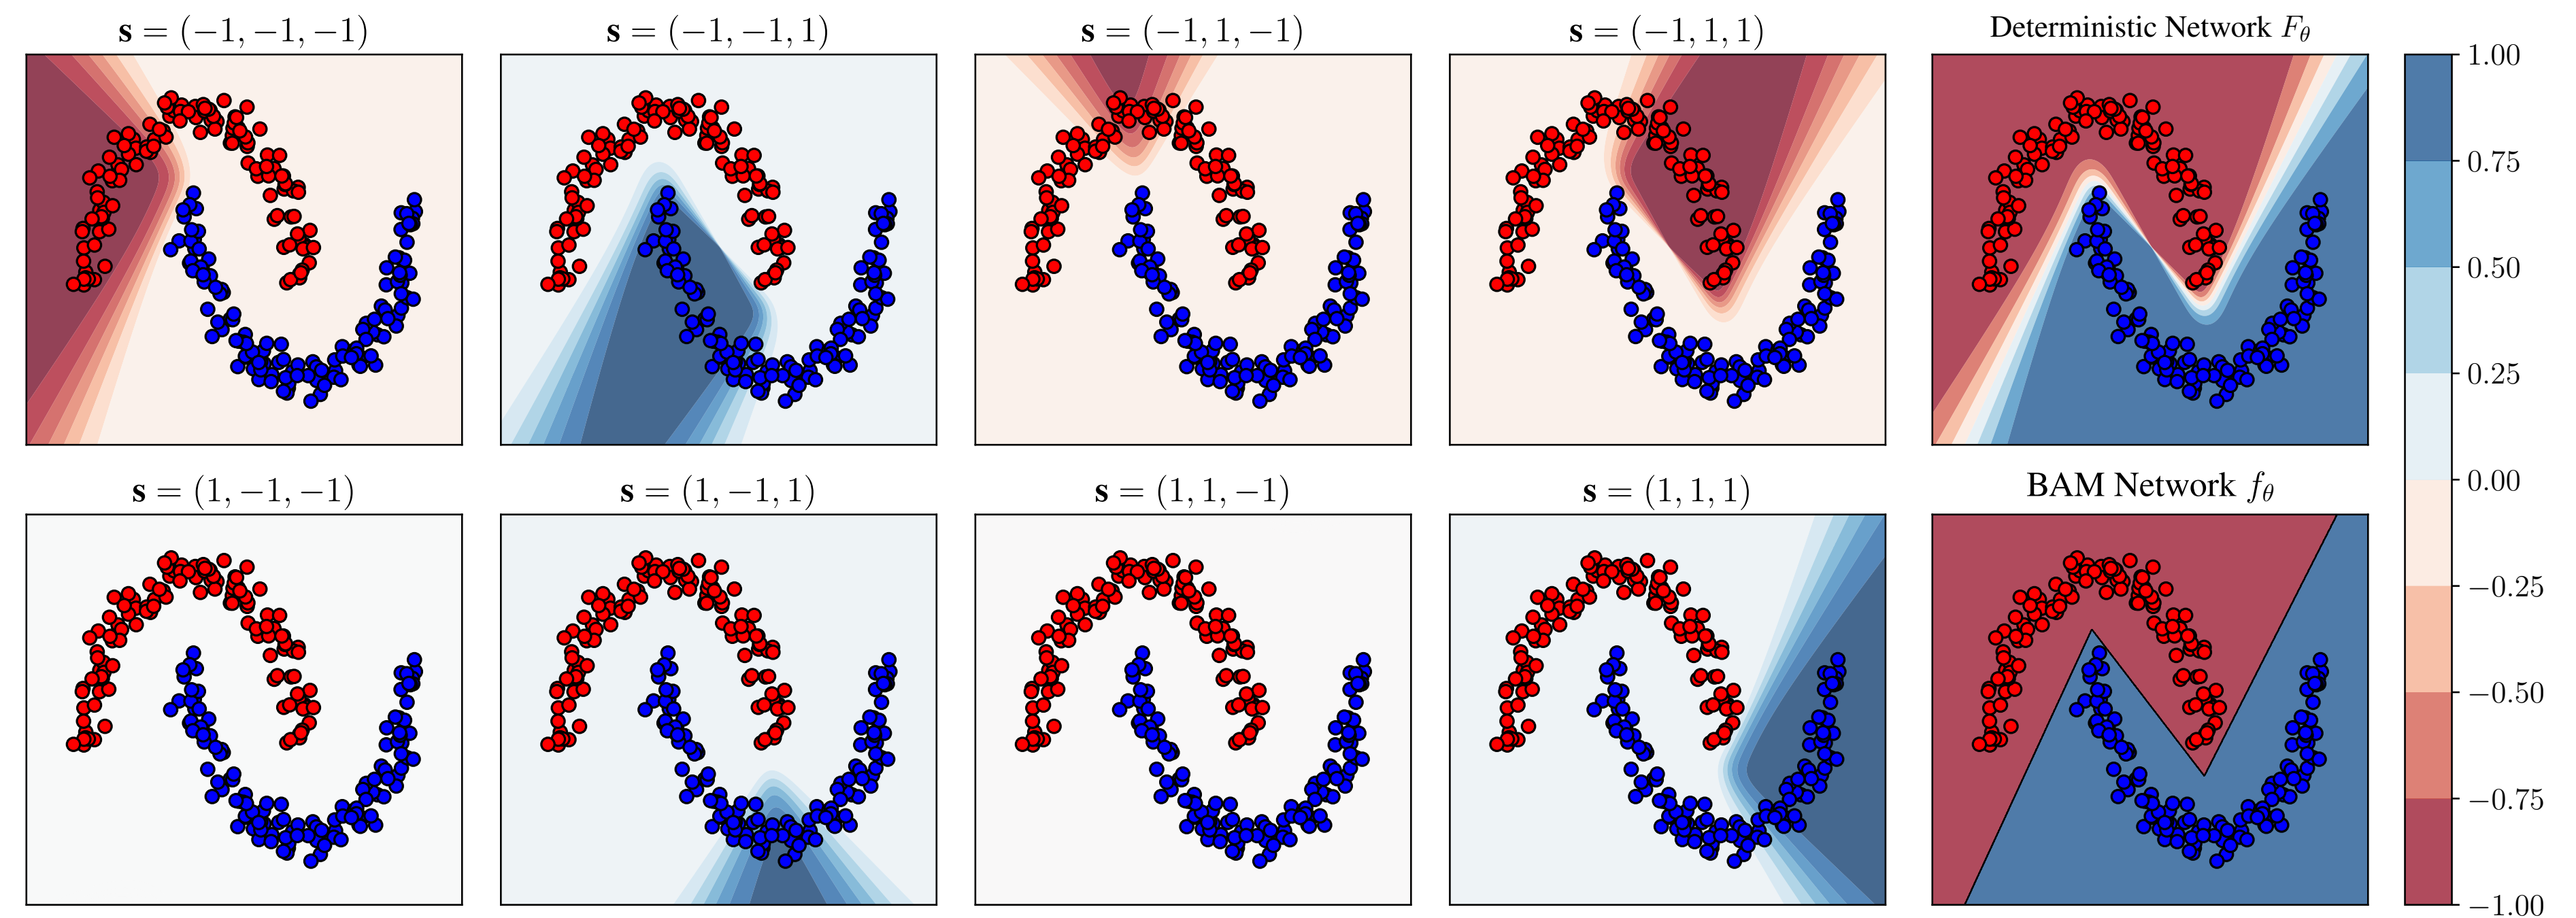


FIGURE 1: Illustration of the proposed method for a one hidden layer network of size  $d_1=3$ , interpreted as a majority vote over 8 binary representations  $\mathbf{s} \in \{-1, 1\}^3$ . For each  $\mathbf{s}$ , a plot shows the values of  $F_{\mathbf{w}_2}(\mathbf{s}) \Pr(\mathbf{s} | \mathbf{x}, \mathbf{W}_1)$ .

## Stochastic Approximation

### Prediction.

$$F_\theta(\mathbf{x}) = \sum_{\mathbf{s} \in \{-1, 1\}^{d_1}} F_{\mathbf{w}_2}(\mathbf{s}) \Pr(\mathbf{s} | \mathbf{x}, \mathbf{W}_1)$$

Hidden layer partial derivatives, with  $\text{erf}'(x) := \frac{2}{\sqrt{\pi}} e^{-x^2}$ .

$$\frac{\partial}{\partial \mathbf{w}_1^k} F_\theta(\mathbf{x}) = \frac{\mathbf{x}}{2^{\frac{3}{2}} \|\mathbf{x}\|} \text{erf}'\left(\frac{\mathbf{w}_1^k \cdot \mathbf{x}}{\sqrt{2} \|\mathbf{x}\|}\right) \sum_{\mathbf{s} \in \{-1, 1\}^{d_1}} s_k F_{\mathbf{w}_2}(\mathbf{s}) \left[ \frac{\Pr(\mathbf{s} | \mathbf{x}, \mathbf{W}_1)}{\Pr(s_k | \mathbf{x}, \mathbf{w}_1^k)} \right]$$

### Monte Carlo sampling

We generate  $T$  random binary vectors  $\{\mathbf{s}^t\}_{t=1}^T$  according to  $\Pr(\mathbf{s} | \mathbf{x}, \mathbf{W}_1)$ .

$$F_\theta(\mathbf{x}) \approx \frac{1}{T} \sum_{t=1}^T F_{\mathbf{w}_2}(\mathbf{s}^t) \quad \frac{\partial}{\partial \mathbf{w}_1^k} F_\theta(\mathbf{x}) \approx \frac{\mathbf{x}}{2^{\frac{3}{2}} \|\mathbf{x}\|} \text{erf}'\left(\frac{\mathbf{w}_1^k \cdot \mathbf{x}}{\sqrt{2} \|\mathbf{x}\|}\right) \frac{1}{T} \sum_{t=1}^T \frac{s_k^t}{\Pr(s_k^t | \mathbf{x}, \mathbf{w}_1^k)} F_{\mathbf{w}_2}(\mathbf{s}^t)$$

This turns out to be a variant of the REINFORCE algorithm (Williams, 1992).

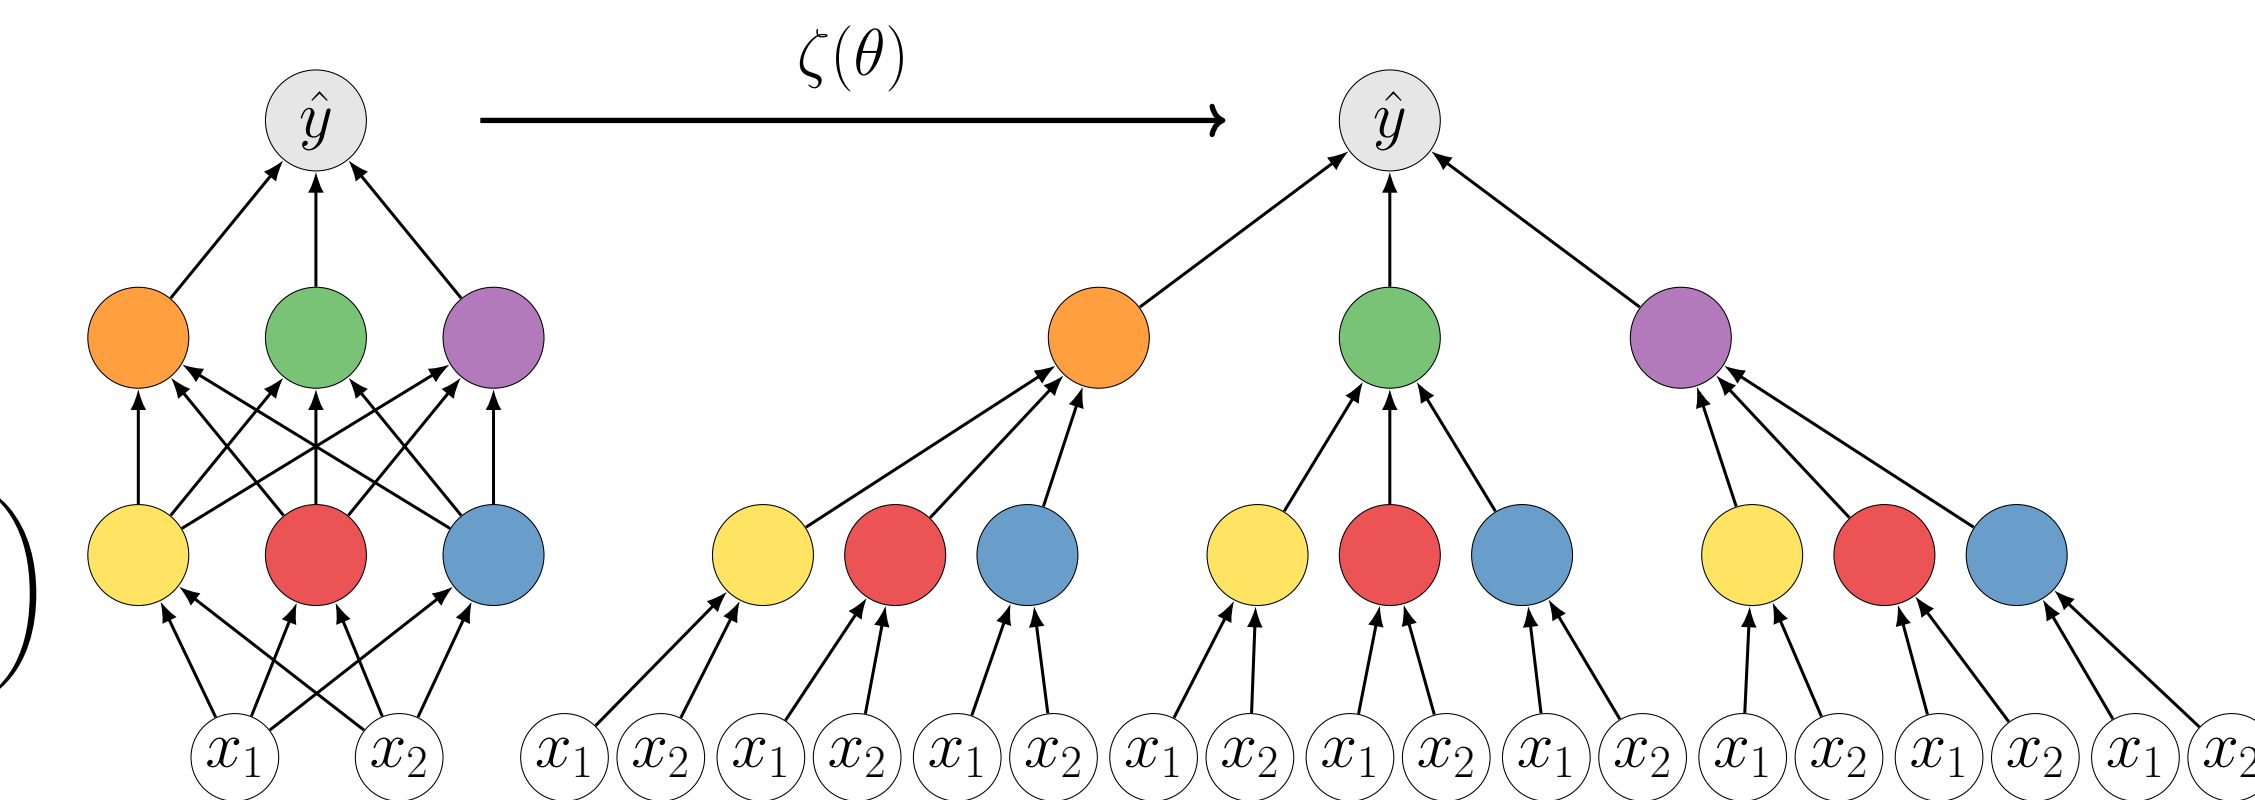
## Deep Learning (PBGNet)

To enable a layer-by-layer computation of the prediction function, we want the neurons of a given layer to be independent of each other. This is achieved with the tree architecture mapping function  $\zeta(\theta)$  applied on a multilayer network.

**Recursive definition.**  $F_k^{(j)}$  denotes the output of the  $j^{\text{th}}$  neuron of the  $k^{\text{th}}$  hidden layer :

$$F_1^{(j)}(\mathbf{x}) = \text{erf}\left(\frac{\mathbf{w}_1^j \cdot \mathbf{x}}{\sqrt{2} \|\mathbf{x}\|}\right) \quad \leftarrow \text{first hidden layer}$$

$$F_{k+1}^{(j)}(\mathbf{x}) = \sum_{\mathbf{s} \in \{-1, 1\}^{d_k}} \text{erf}\left(\frac{\mathbf{w}_{k+1}^j \cdot \mathbf{s}}{\sqrt{2} d_k}\right) \prod_{i=1}^{d_k} \left( \frac{1}{2} + \frac{1}{2} s_i \times F_k^{(i)}(\mathbf{x}) \right)$$



### Kullback-Leibler regularization

$$\text{KL}(Q_{\zeta(\theta)} \| P_{\zeta(\mu)}) = \frac{1}{2} \left( \|\mathbf{w}_L - \mathbf{u}_L\|^2 + \sum_{k=1}^{L-1} d_{k+1}^\dagger \|\mathbf{W}_k - \mathbf{U}_k\|_F^2 \right) \text{ for } d_k^\dagger := \prod_{i=k}^L d_i.$$

## Experiment

We perform model selection using a validation set for a MLP with tanh activations, and using the PAC-Bayes bound for our PBGnet and PBGnet\_pre algorithms. PBGnet\_t and PBGnet\_t-bnd are intermediate variants.




FIGURE 2: Experiment results for the considered models on the binary classification datasets. PAC-Bayesian bounds hold with probability 0.95.